




Clean (Unattacked) Accuracy — 7 Models




Axis 1 — Gradient Attacks per Model
PGD/AutoAttack/Square drive every model to ~0% (no gradient masking)




Axis 1 — Accuracy Drop (clean – robust), % · darker = worse




Axis 2 — Clean vs Transfer-Robust
(ResNet-50 surrogate; R50 row = white-box self-transfer)




Transfer Fooling Rate
CNN→Transformer transfers poorly (ViT most resistant)




★ KEY FINDING — CLIP TASR = 34.2%
(pure-classifier mean $\approx 0.1\%$)




Typographic Accuracy Drop
CLIP -25.3 pp vs pure-mean -1.6 pp (16.3×)




Axis 4 — ImageNet-C-style Corruptions: Robust Accuracy (%), severity 3
green = robust, red = fragile · italic columns = NumPy approximations



Axis 4 — Mean Corruption Robustness per Model



Robustness Radar — 7 Baselines (higher = better)
PGD/AutoAttack omitted: all models ≈ 0 there



Axis 1 — Gradient Attacks: Full Accuracy Table (7 Models)

green = robust, red = fragile

	Clean	FGSM	PGD-20	AutoAttack	Square
	0.781 (78.1%)	0.408 (40.8%)	0.004 (0.4%)	0.000 (0.0%)	0.010 (1.0%)
	0.687 (68.7%)	0.011 (1.1%)	0.001 (0.1%)	0.000 (0.0%)	0.000 (0.0%)
	0.787 (78.7%)	0.292 (29.2%)	0.000 (0.0%)	0.000 (0.0%)	0.010 (1.0%)
	0.793 (79.3%)	0.244 (24.4%)	0.000 (0.0%)	0.000 (0.0%)	0.000 (0.0%)
	0.775 (77.5%)	0.191 (19.1%)	0.000 (0.0%)	0.000 (0.0%)	0.000 (0.0%)
-B0	0.743 (74.3%)	0.100 (10.0%)	0.000 (0.0%)	0.000 (0.0%)	0.000 (0.0%)
6	0.621 (62.1%)	0.058 (5.8%)	0.000 (0.0%)	0.000 (0.0%)	0.000 (0.0%)

Axis 2 — Transfer Attacks: Full Table (surrogate = ResNet-50 PGD-20)

ResNet-50 row = white-box self-transfer (validates harness: 0.003 ≈ PGD 0.004)

	White-box?	Clean	Transfer-Robust	Drop	Fooling Rate
	yes	0.781 (78.1%)	0.003 (0.3%)	0.778 (77.8%)	0.996 (99.6%)
	no	0.695 (69.5%)	0.478 (47.8%)	0.217 (21.7%)	0.327 (32.7%)
	no	0.783 (78.3%)	0.638 (63.8%)	0.145 (14.5%)	0.198 (19.8%)
	no	0.796 (79.6%)	0.747 (74.7%)	0.049 (4.9%)	0.079 (7.9%)
	no	0.772 (77.2%)	0.661 (66.1%)	0.111 (11.1%)	0.161 (16.1%)
-B0	no	0.748 (74.8%)	0.623 (62.3%)	0.125 (12.5%)	0.186 (18.6%)
6	no	0.628 (62.8%)	0.532 (53.2%)	0.096 (9.6%)	0.205 (20.5%)

Fooling rate = fraction of originally-correct images flipped. CNN→Transformer transfers poorly (ViT-B/16 most resistant).

Axis 3 — Typographic Attack: Full Table

★ CLIP TASR = 34.2% vs pure-classifier mean ≈ 0.1% (CLIP TASR cell highlighted orange)

	Clean	Robust	Drop	Fooling Rate	TASR (targeted)
	0.781 (78.1%)	0.768 (76.8%)	0.013 (1.3%)	0.036 (3.6%)	0.001 (0.1%)
	0.687 (68.7%)	0.663 (66.3%)	0.024 (2.4%)	0.076 (7.6%)	0.000 (0.0%)
	0.787 (78.7%)	0.777 (77.7%)	0.010 (1.0%)	0.032 (3.2%)	0.000 (0.0%)
	0.793 (79.3%)	0.773 (77.3%)	0.020 (2.0%)	0.039 (3.9%)	0.001 (0.1%)
	0.775 (77.5%)	0.770 (77.0%)	0.005 (0.5%)	0.032 (3.2%)	0.002 (0.2%)
-B0	0.743 (74.3%)	0.722 (72.2%)	0.021 (2.1%)	0.055 (5.5%)	0.001 (0.1%)
6	0.621 (62.1%)	0.368 (36.8%)	0.253 (25.3%)	0.452 (45.2%)	0.342 (34.2%)

TASR = fraction of images predicted as the exact wrong class written on the sticker. Pure classifiers cannot read text → TASR ≈ 0. CLIP is language-grounded → TASR = 34.2%.

Axis 4 — Corruptions: Full Numeric Table (robust accuracy per cell, severity 3)

Color: green = robust, red = fragile

	gaussian noise	shot noise	impulse noise	defocus blur	glass blur	motion blur	zoom blur	snow	frost	fog	brightness	contrast	elastic transform	pixelate	jpeg compression	Mean
	0.487	0.457	0.358	0.447	0.088	0.129	0.393	0.664	0.618	0.384	0.721	0.700	0.771	0.594	0.660	0.498
	0.150	0.143	0.091	0.190	0.056	0.045	0.242	0.511	0.420	0.165	0.563	0.313	0.647	0.285	0.475	0.286
r	0.567	0.554	0.565	0.489	0.130	0.160	0.417	0.730	0.656	0.387	0.750	0.712	0.769	0.643	0.701	0.549
	0.636	0.637	0.515	0.522	0.258	0.216	0.463	0.730	0.677	0.543	0.743	0.717	0.773	0.686	0.695	0.587
	0.592	0.566	0.563	0.490	0.134	0.181	0.408	0.727	0.684	0.549	0.738	0.728	0.773	0.583	0.671	0.559
t-B0	0.432	0.436	0.425	0.353	0.093	0.110	0.337	0.609	0.536	0.226	0.695	0.526	0.725	0.585	0.647	0.449
16	0.403	0.400	0.368	0.396	0.147	0.141	0.316	0.533	0.511	0.355	0.579	0.543	0.596	0.481	0.488	0.417

Master Comparison — All 7 Models × All 4 Axes												
	G:Clean	G:FGSM	G:PGD	G:AA	G:Sq	T:Rob	T:Drop	T:Fool	Ty:Rob	Ty:Drop	Ty:TASR	C:Mean
	0.781	0.408	0.004	0.000	0.010	0.003	0.778	0.996	0.768	0.013	0.001	0.498
	0.687	0.011	0.001	0.000	0.000	0.478	0.217	0.327	0.663	0.024	0.000	0.286
	0.787	0.292	0.000	0.000	0.010	0.638	0.145	0.198	0.777	0.010	0.000	0.549
	0.793	0.244	0.000	0.000	0.000	0.747	0.049	0.079	0.773	0.020	0.001	0.587
	0.775	0.191	0.000	0.000	0.000	0.661	0.111	0.161	0.770	0.005	0.002	0.559
-B0	0.743	0.100	0.000	0.000	0.000	0.623	0.125	0.186	0.722	0.021	0.001	0.449
6	0.621	0.058	0.000	0.000	0.000	0.532	0.096	0.205	0.368	0.253	0.342	0.417

G: Gradient (Clean/FGSM/PGD/AA/Square)

T: Transfer (Robust/Drop/Fool)

Ty: Typographic (Robust/Drop/TASR)

C: Corruption mean (15 types)